

- ❑ **Part 1: Symbolic AI representations and tasks**
 - ❖ (Sub)symbolic AI, and representations
 - ❖ Welcome/Purpose of this tutorial
 - ❖ Question Answering on Knowledge Graphs
- ❑ **Part 2: Text-to-symbols and Ranking**
 - ❖ Neural Text and Graph Representations
 - ❖ Text-Symbol Alignment and Semantic Annotations
 - ❖ Entity Representations and Entity Ranking
- ❑ **Part 3: Neuro-symbolic representations for Reasoning**
 - ❖ Reasoning about Relevance
 - ❖ Neuro Pseudo-Relevance Feedback with Explainability
- ❑ **Part 4: Applications for Neuro-symbolic approaches**
 - ❖ Use Case: Knowledge Discovery
 - ❖ Use Case: Task-based Assistance
- ❑ **Panel & Discussion**

Material: <https://github.com/laura-dietz/neurosymbolic-representations-for-IR/>

Entity Representations and Entity Ranking

DR. SHUBHAM CHATTERJEE


RESEARCH ASSOCIATE

DEPARTMENT OF COMPUTING SCIENCE

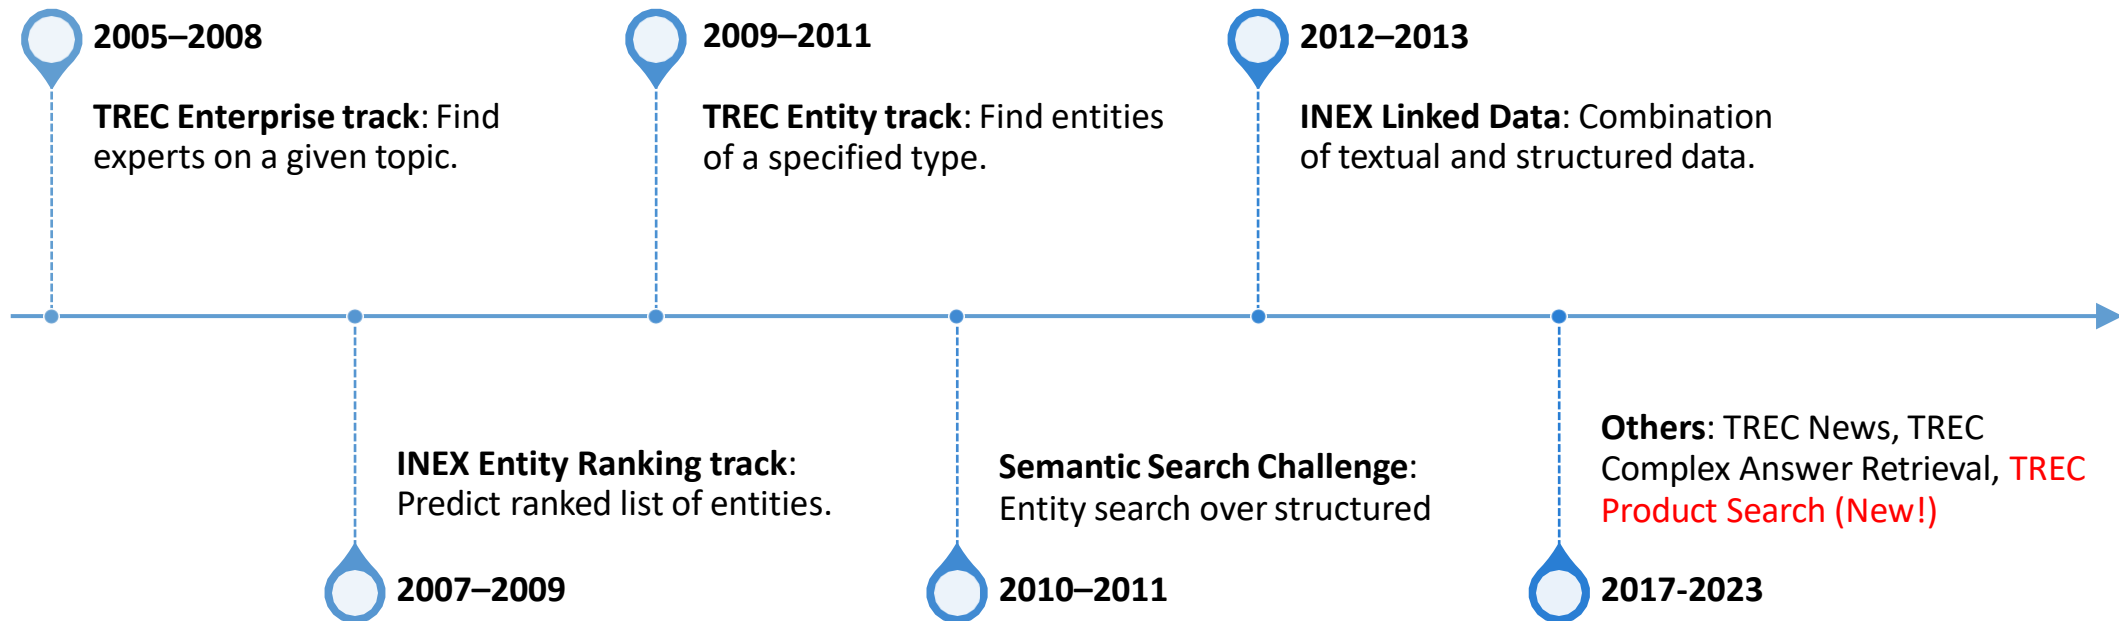
UNIVERSITY OF GLASGOW, UK

Task: Entity Ranking

Given a query and a KG, retrieve entities that are relevant to the query ordered by the relevance of each entity to the query.



Entity Ranking: Over the Years




How Do Systems Usually Retrieve Entities?

Probabilistic Unstructured Models

Consider an entity as a single “document”.

- ☐ Create a search index of all entities.
- ☐ Match the query against the entity representation.

One LM for whole entity



Probabilistic Unstructured Models

Consider an entity as a single “document”.

- ☐ Create a search index of all entities.
- ☐ Match the query against the entity representation.

One LM for whole entity


- ☐ Probability distribution over sequences of words.
- ☐ Generate probabilities by training on text corpora in one or many languages.

Fielded Models

Consider an entity as a “document” with multiple fields.

- ☐ Create a search index of all entities.
- ☐ Match the query against the entity representation.

Multiple LMs for different fields

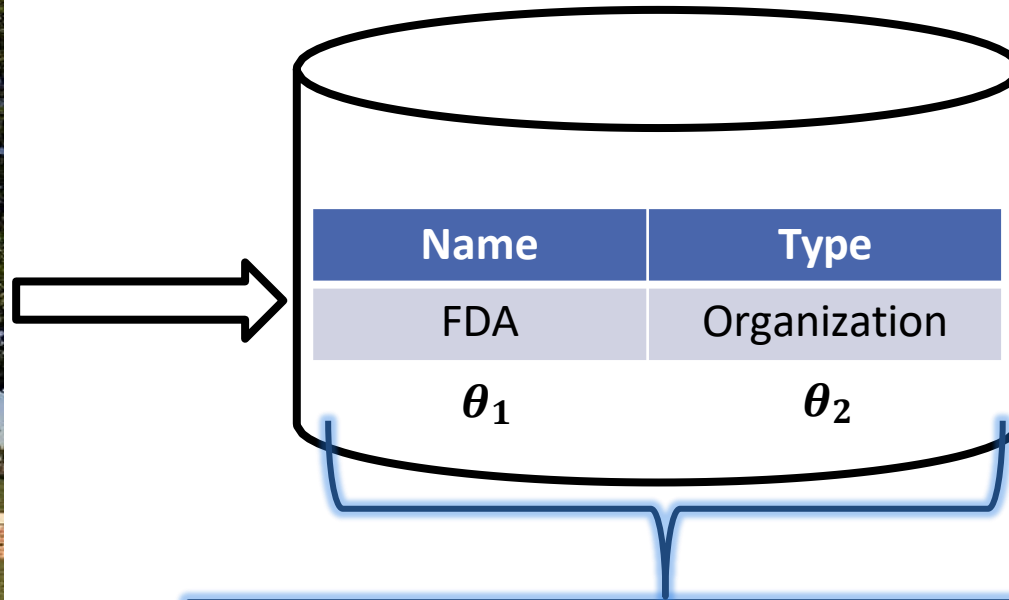


Fielded Models: Multiple LM for Different Fields

Represent entities as “documents” with different fields.



Entity: FDA



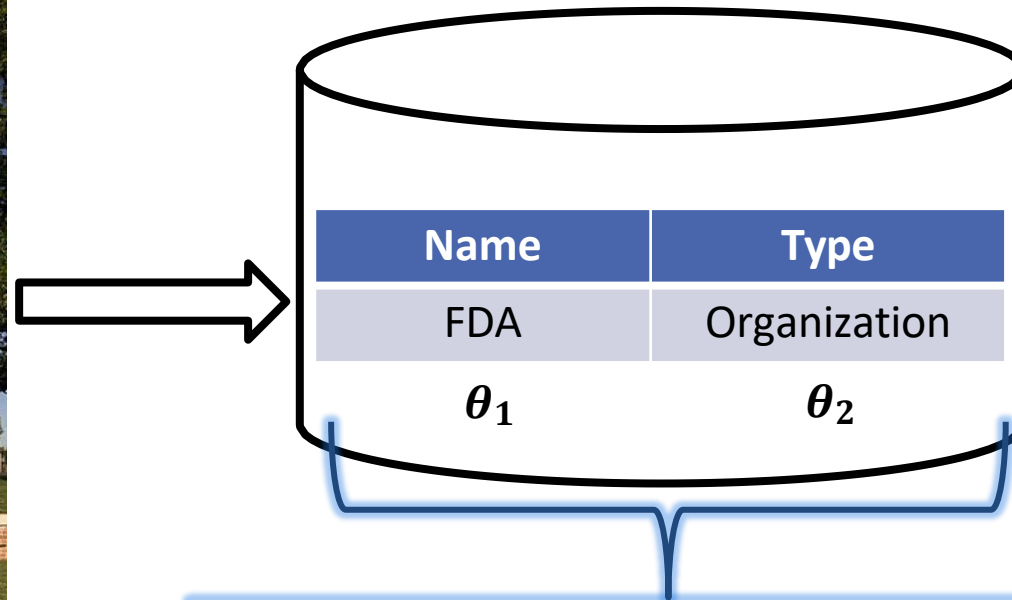
Different language models for different fields.

Fielded Models: Multiple LM for Different Fields



Entity: FDA

Represent entities as “documents” with different fields.



Combine using a linear mixture with field weights

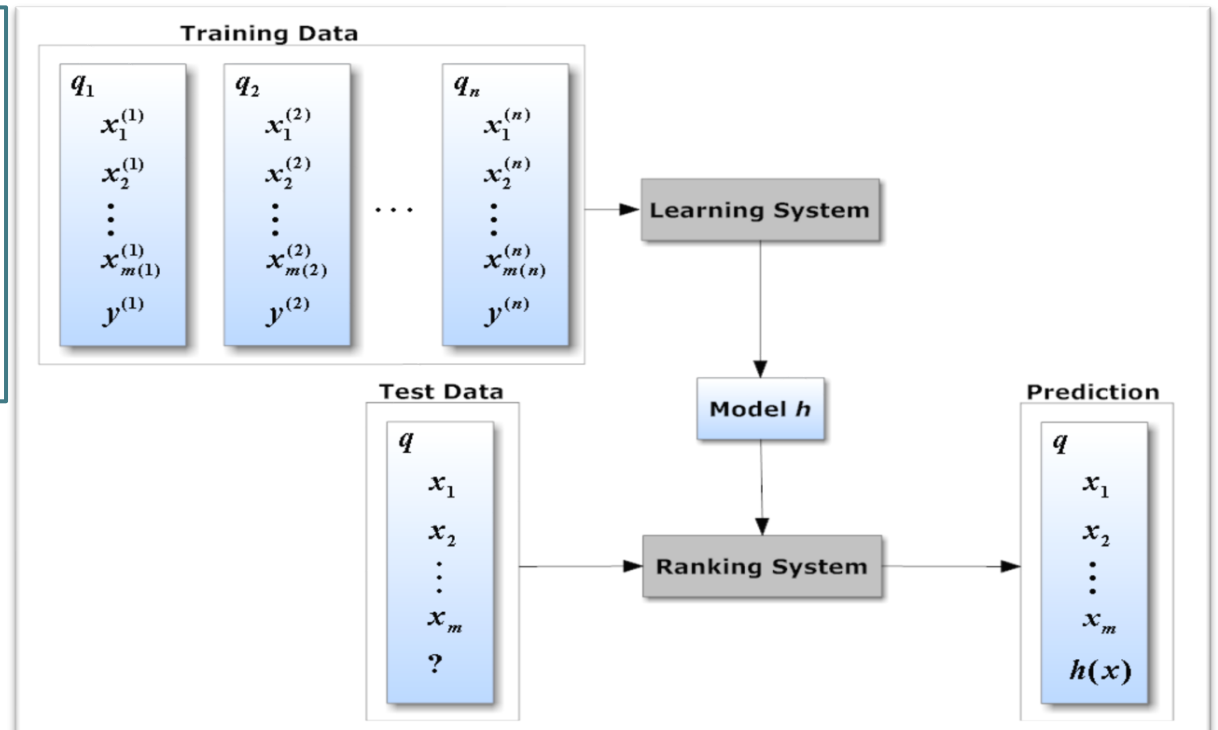
Use Coordinate Ascent

Learning-to-Rank Models

Use feature vectors of (query, entity) pairs to train a ML model.

Features:

- ❑ Entity retrieval from a KB using BM25+RM3.
- ❑ Whether the candidate entity is contained in the query entities.
- ❑ Normalized Levenshtein Distance between the query and the mention.



Picture credit: <http://web.ist.utl.pt/~catarina.p.moreira/coursera.html>

Type-Aware Entity Retrieval


- ❑ Uses entity types from a type taxonomy (e.g., Wikipedia categories).
- ❑ Query enriched with types of entities in the query (called target types)
- ❑ Example [Balog et al., 2011]
 - Learn a probability distribution over the query and entity types.
 - Similarity = KL divergence between two distributions

How Do Systems Usually Retrieve Entities?

- ❑ Create a search index of all entities
- ❑ Match the query against the **representation** of the entity.

Non-neural methods: Sparse

Neural methods: Dense



Neural Entity Ranking

GEEER: Using Wikipedia2Vec for Entity Ranking

- Re-ranks entities using Wikipedia2Vec.
- Shows that Wikipedia2Vec is useful for entity ranking.
- **General Idea:** Relevant entities for a given query are situated close (in graph embedding space) to the query entities identified by the entity linker.

□ Method:

- ❖ Compute the embedding-based score for an entity:

$$Score_{emb}(E, Q) = \sum_{e \in Q} C(e) \cdot \cos(\vec{E}, \vec{e})$$

- ❖ Final Score = interpolation of the embedding-based and retrieval scores.

$$Score_{final}(E, Q) = \lambda \cdot Score_{emb}(E, Q) + (1 - \lambda) \cdot Score_{ret}(E, Q)$$

GEEER: Using Wikipedia2Vec for Entity Ranking

- Re-ranks entities using Wikipedia2Vec.
- Shows that Wikipedia2Vec is useful for entity ranking.
- **General Idea:** Relevant entities for a given query are situated close (in graph embedding space) to the query entities identified by the entity linker.
- **Method:**

- ❖ Compute the embedding-based score for an entity:

$$Score_{emb}(E, Q) = \sum_{e \in Q} C(e) \cdot \cos(\vec{E}, \vec{e})$$

Confidence score of
entity linker.


- ❖ Final Score = interpolation of the embedding-based and retrieval scores.

$$Score_{final}(E, Q) = \lambda \cdot Score_{emb}(E, Q) + (1 - \lambda) \cdot Score_{ret}(E, Q)$$

Use LTR optimized for
NDCG@100.


KEWER: Knowledge graph Entity and Word Embedding for Retrieval

- Joint embedding of entities and words (like Wikipedia2Vec).
- **Uses:** Entities, attributes, categories, KG structure.
- **Method:** Random walks over a KG.
 - ❖ Start from each entity.
 - ❖ Repeatedly follow directed edges
- **Entity Ranking:** Cosine similarity between query and entity embeddings.



Graph Embeddings: Issues for IR

- ❑ IR: Considers explicit query.
- ❑ Current graph embedding methods: Not seen the query during training (**query-agnostic!**)



Picture Credit: Mengjia Xu. *Understanding Graph Embedding Methods and their Applications*. 2020.

How About We Use BERT?




- ☐ BERT shown to be useful for document retrieval.
- ☐ **Question: Can we use BERT for entity ranking?**

One Idea: E-BERT: Injecting Entity information into BERT

- ❑ Aligns Wikipedia2Vec entity vectors with BERT's native word piece vectors.
- ❑ **Issue:** BERT's word piece dictionary does not contain any entities!
- ❑ **Solution:** Use common words in the vocabulary of BERT and Wikipedia2Vec.
- ❑ **Idea:** Wikipedia2Vec embeds words and entities into the same vector space → Alignment learnt using words can also be applied to entities.

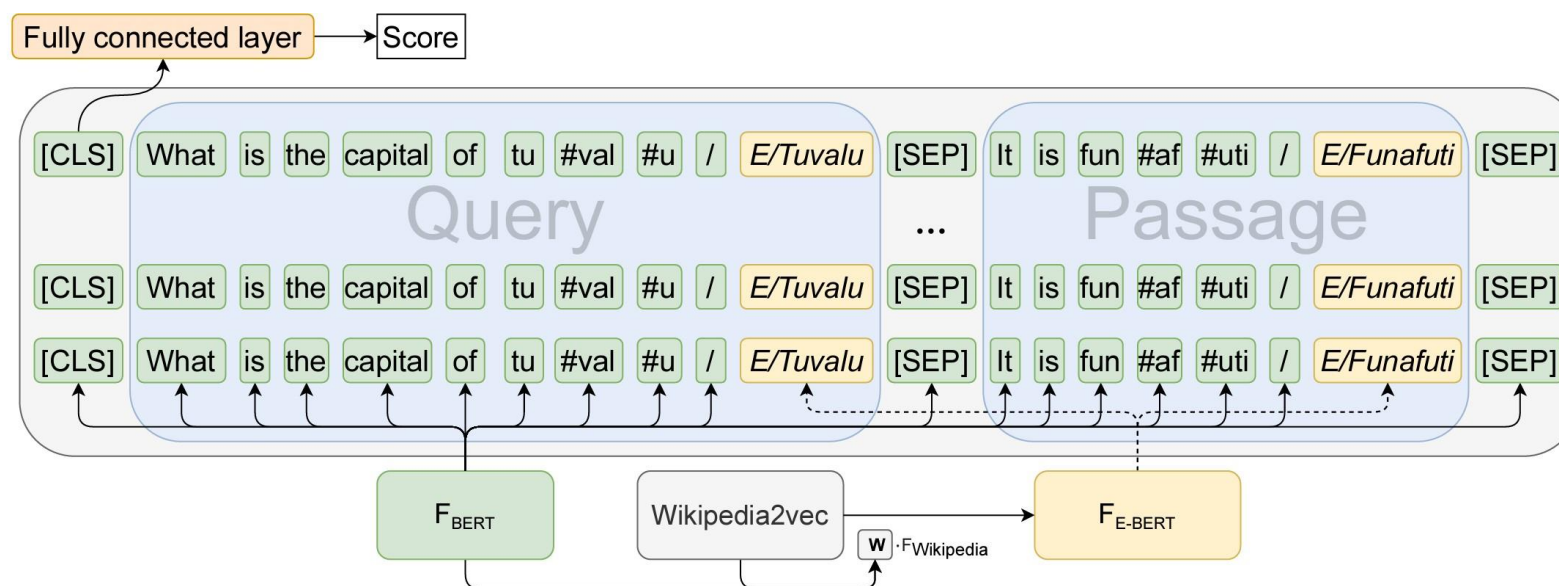
How Well Does BERT Understand Entities?



Results for GEEER with ERNIE and E-BERT. DBpedia-Entity v2.

EM-BERT: Entity-enriched BERT for Entity Ranking

- Aligns Wikipedia2Vec entity vectors with BERT's native word piece vectors (as in E-BERT).
- Entity vectors used with entity mentions.
- Fine-tuning: First MS MARCO passages, then DBpedia-Entity v2.



Picture Credit: Gerritse et al. Entity-aware Transformers for Entity Search. 2022.

Figure 1: Illustration of the EM-BERT model. Entity annotated query and documents are tokenized and mapped to their corresponding vector representations using F_{BERT} and F_{E-BERT} functions.

EM-BERT: What do we learn?

- ❑ Substantial improvements over SOTA for entity ranking!
- ❑ Helps:
 - ❖ Complex natural language queries,
 - ❖ List search queries, and
 - ❖ Queries containing tail entities

One Idea: Injecting Entity information into BERT

Another Idea: Can we utilize existing knowledge in BERT?

- ❑ BERT has already seen a lot of the world (from books and Wikipedia).
- ❑ BERT can probably infer the connection between the query and entity from a term-based entity description.
- ❑ Term-based entity description = Introductory Wikipedia paragraph (most often).


Query: Genetically Modified Organism

Relevant Entity: Food and Drug Administration

Food and Drug Administration

Lead Text

The United States **Food and Drug Administration (FDA or USFDA)** is a [federal agency](#) of the [Department of Health and Human Services](#). The FDA is responsible for protecting and promoting [public health](#) through the control and supervision of [food safety](#), [tobacco products](#), [dietary supplements](#), [prescription](#) and [over-the-counter pharmaceutical drugs](#) (medications), [vaccines](#), [biopharmaceuticals](#), [blood transfusions](#), [medical devices](#), [electromagnetic radiation emitting devices \(ERED\)](#), [cosmetics](#), [animal foods & feed](#)^[3] and [veterinary products](#).




- Lead Text = Static entity description
- No knowledge of the query!
- Static entity embeddings.

Query: Genetically Modified Organism

Relevant Entity: Food and Drug Administration

FDA regulates most human and animal food, including GMO foods. In doing so, FDA makes sure that foods that are GMOs or have GMO ingredients meet the same strict safety standards as all other foods. FDA sets and enforces food safety standards [...]

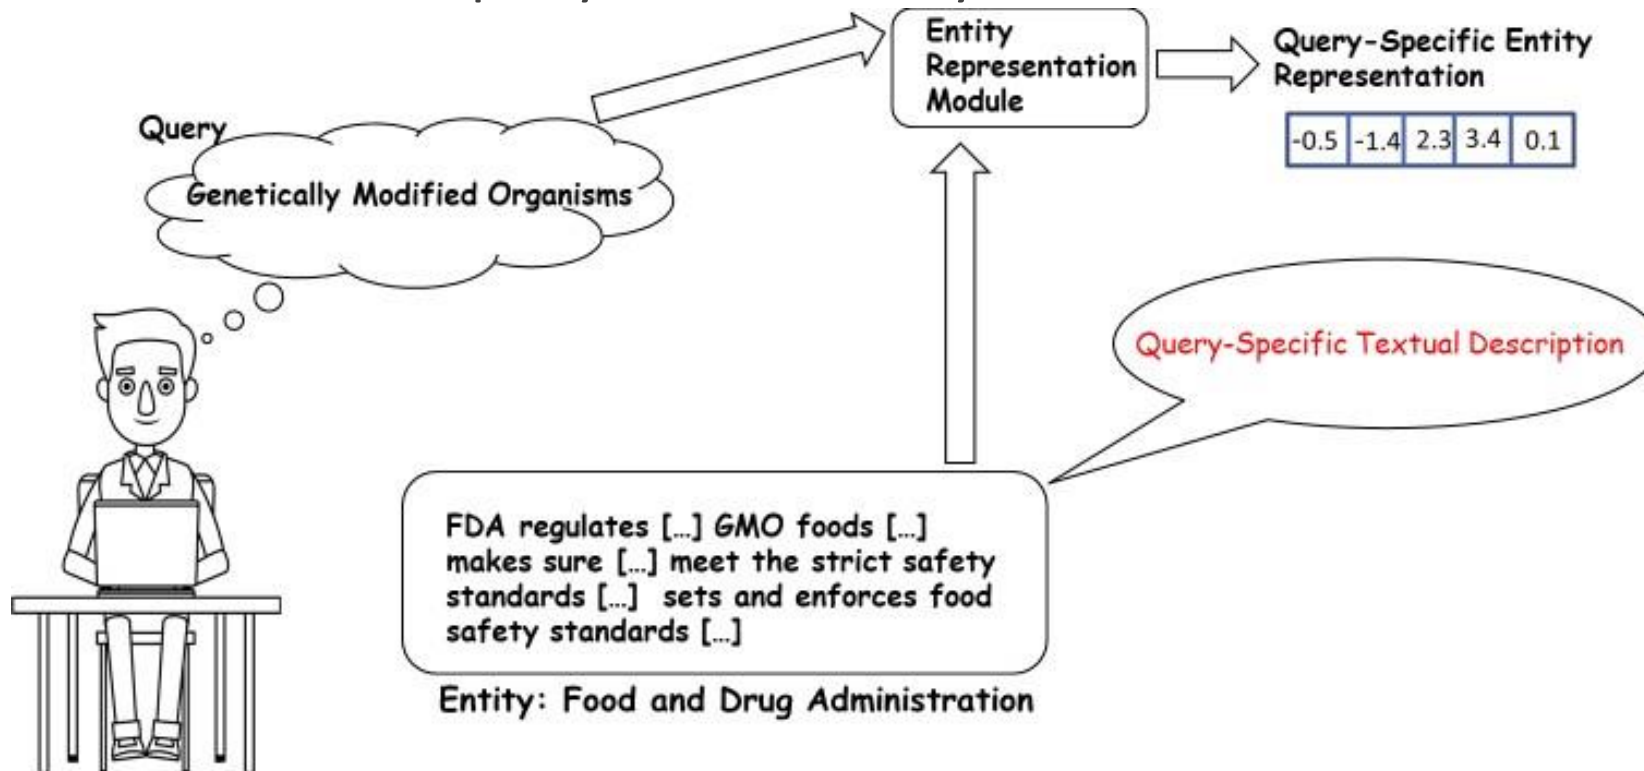
**Clarifies the connection
between the query and entity.**



What if we had this text instead of lead text?

BERT-ER: Query-Specific Entity Descriptions

Query-specific entity descriptions → descriptions that mention relevant connections between the query and the entity.



Food and Drug Administration

Chatterjee and Dietz., 2022

The United States **Food and Drug Administration** (FDA or **USFDA**) is a federal agency of the Department of Health and Human Services. The FDA is responsible for protecting and promoting public health through the control and supervision of food safety, tobacco products, dietary supplements, prescription and over-the-counter pharmaceutical drugs (medications), vaccines, biopharmaceuticals, blood transfusions, medical devices, electromagnetic radiation emitting devices (ERED), cosmetics, animal foods & feed^[3] and veterinary products.

- 1 [Organizational structure](#)
- 2 [Location](#)
- 3 [Scope and funding](#)
- 4 [Regulatory programs](#)
- 5 [Science and research programs](#)
- 6 [Data management](#)
- 7 [History](#)

Aspects = Top-Level Sections

Query-Specific Entity Descriptions: Alternative 1

❑ Using Wikipedia: Top-Level Sections

❑ Identify relevant top-level sections from the Wikipedia page. (*Why? —because the lead text is does not elaborate the relevance!*)


❑ Use catalog of top-level sections (aspects) from Ramsdell et al., 2020.

❑ **Downside: Wikipedia articles often do not contain all relevant information!**

Query-Specific Entity Descriptions: Alternative 2

Chatterjee and Dietz., 2022

❑ Using Paragraph Collection: PRF Passages




❑ **Downside:** *Entity may not be central to the discussion in the text!*

Query-Specific Entity Descriptions: Alternative 3

Chatterjee and Dietz., 2022

❑ Using Paragraph Collection: Entity-Support Passages



❑ Re-rank these documents.

❑ Two criteria:

1. How many relevant connections between query and entity?
2. Are the relevant connections central to the discussion in the text?

Entity-Support Passages:
Chatterjee and Dietz, ICTIR 2019.

BERT-ER: Alternative Approaches for Query-Specific Entity Descriptions

Chatterjee and Dietz, SIGIR 2022

Query: Genetically Modified Organism

Entity: Food and Drug Administration

Entity-Linked Corpus

Lead (Baseline)

The Food and Drug Administration (FDA) is a federal agency of the Department of Health and Human Services. ...

Aspect

Location

Regulatory programs
... The programs for safety regulation vary widely by the type of products, its potential risks, and the regulatory power ...

Lead

Aspect

Aspect

Aspect

Food and Drug Administration

From Wikipedia, the free encyclopedia

"FDA" redirects here. For other uses, see FDA (disambiguation).
Not to be confused with the Drug Enforcement Administration.

The United States **Food and Drug Administration** (**FDA** or **USFDA**) is a federal agency of the Department of Health and Human Services. The FDA is responsible for protecting and promoting public health through the control and supervision of food safety, tobacco products, dietary supplements, prescription and over-the-counter pharmaceutical drugs (medications), vaccines, biopharmaceuticals, blood transfusions, medical devices, electromagnetic radiation emitting devices (ERED), cosmetics, animal foods & feed^[2] and veterinary products.

Organizational structure [\[edit\]](#)
Department of Health and Human Services

Location [\[edit\]](#)

Headquarters [\[edit\]](#)
FDA headquarters facilities are currently located in [Montgomery County and Prince George's County, Maryland](#).^[3]

Other locations [\[edit\]](#)
The FDA has a number of field offices across the United States, in addition to international locations in China, India, Europe, the Middle East, and Latin America.^[4]

Regulatory programs [\[edit\]](#)

Emergency approvals (EUA) [\[edit\]](#)

Emergency Use Authorization (EUA) is a mechanism that was created to facilitate the availability and use of medical countermeasures, including vaccines and personal protective equipment, during public health emergencies such as the Zika virus epidemic, the Ebola virus epidemic and the COVID-19 pandemic.^[5]

Regulations [\[edit\]](#)

Main article: Regulation of food and dietary supplements by the U.S. Food and Drug Administration

A more recent example of the FDA's international work is their 2018 cooperation with

Food and Drug Administration

U.S. FOOD & DRUG ADMINISTRATION

Agency overview

Formed June 30, 2002, 125 years ago^[1]

Jurisdiction Federal government of the United States

Headquarters White Oak Campus

10903 New Hampshire Avenue

Silver Spring, Maryland 20910

Area served 18,000 (2022)^[2]

Employees 18,000 (2022)^[2]

Agency executives Janet Woodcock (acting Commissioner)

Ann Avenet, Principal Deputy Commissioner

Parent agency Department of Health and Human Services

Child agencies Office of Criminal Investigations

Office of Regulatory Affairs

Website [www.fda.gov](#)



BERT

Query-specific Entity Representations

PRF-Passage

A genetically modified organism (GMO) is [...] . **FDA** regulates most human and animal food, including GMO foods.[..]

Entity-Support Passage

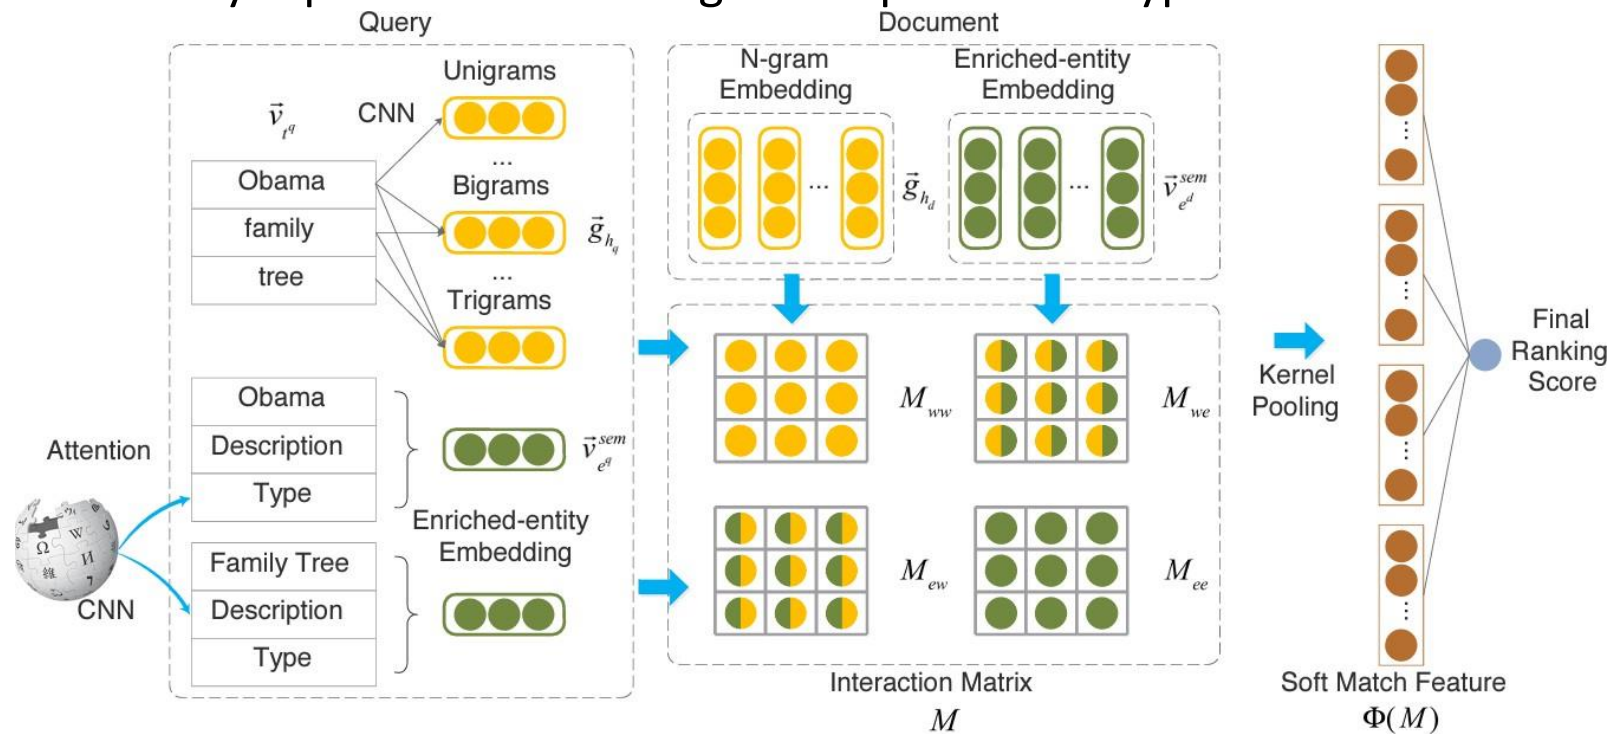
The U.S. **Food and Drug Administration** (FDA) ensures that GMOs are safe for human, plant, and animal health. ...

How Do We Use Entities?

Entity-Duet Neural Ranking Model

Liu et al., 2018

- ❑ Incorporates entities in interaction-based neural ranking models.
- ❑ Learns entity representations using: descriptions and types.



Picture Credits: Liu et al. Entity-Duet Neural Ranking. 2018.

Dense Retrieval With Entity Views

Tran and Yates, 2022

- Enrich query/document representation with entity representations.
- Cluster entities → Entity embeddings using clusters → Clusters act as “views” of the document.

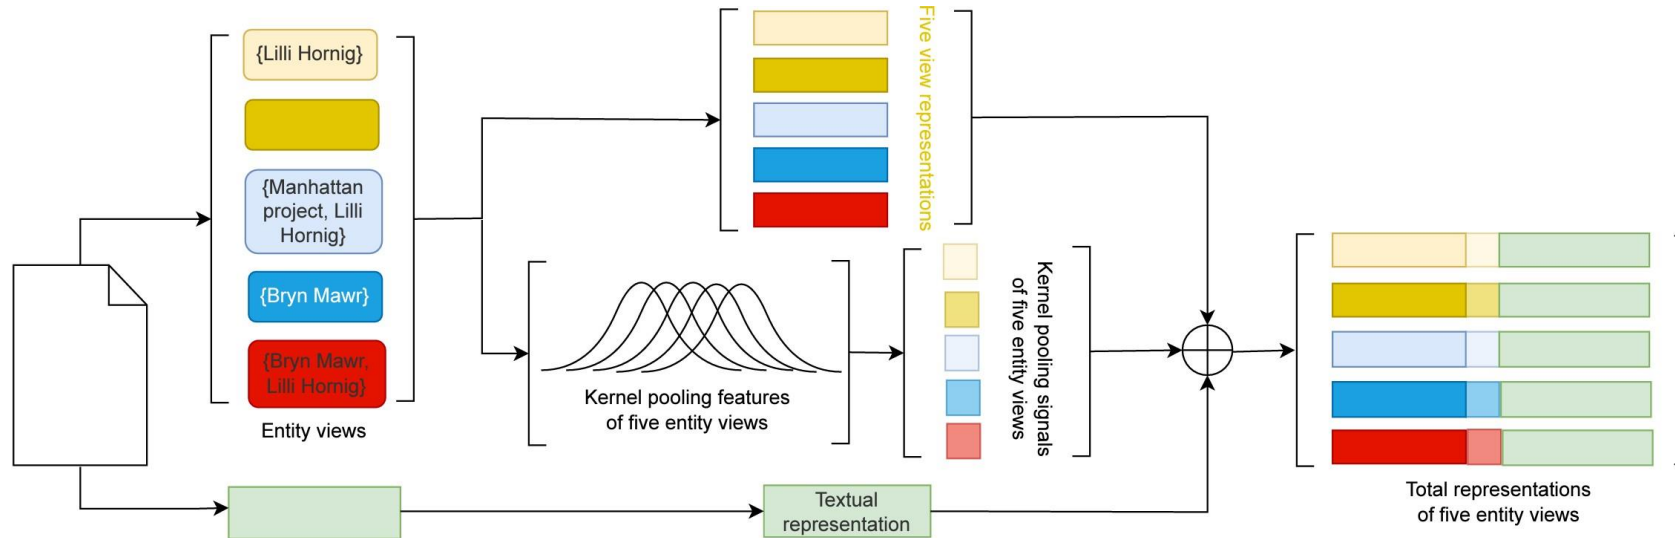


Figure 1: Overview of EVA with multiple representations. Entity clusters such as {Lilli Hornig}, {Manhattan project}, {Manhattan project, Lilli Hornig}, {Bryn Mawr} and {Bryn Mawr, Lilli Hornig} can be understood as different entity views of the passage. EVA generates one total representation for each view, which enriches a textual representation with the entities present.

Picture Credits: Tran and Yates. Dense Retrieval with Entity Views. 2022.