

Knowledge Graphs & Entity Linking

(and their relation to LLMs)

Session 1 of the SIGIR'23 Tutorial
"Neuro-Symbolic Representations for IR"
Taipei, Taiwan, July 23, 2023

Hannah Bast

Chair of Algorithms and Data Structures
University of Freiburg, Germany

Knowledge Graphs 1/8



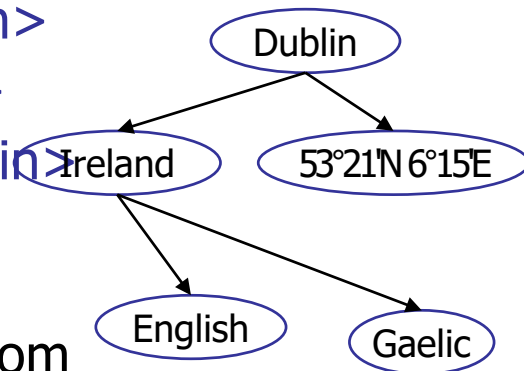
■ Resource Description Framework

- RDF is based on a strikingly simple principle: ALL DATA, whatever it is, is modelled as a **set of triples**

subject predicate object

- For example:

| | | |
|----------|---------------|--------------------|
| <Taipei> | <capital of> | <Taiwan> |
| <Taipei> | <coordinates> | "25°02'N 121°34'E" |
| <Taiwan> | <language> | <Hokkien> |
| <Taiwan> | <language> | <Hakka> |
| <Taiwan> | <language> | <Mandarin> |



Graph view: each distinct subject or object is a **node**, each triple is a **directed edge** from subject to object, with the predicate as label



Knowledge Graphs 2/8

wdt:P17 ~ country
wd:Q865 ~ Taiwan
wdt:P625 ~ coordinates

■ International Resource Identifier (IRI)

- Names have globally unique identifiers, so-called IRIs

An IRI is like an address in a browser (URL), only that Unicode characters like ä, ö, ü, ... are allowed

- The IRI is often not human-readable but ID-like; instead there are dedicated triples for its names and aliases
- For example, the Wikidata IRI for **Taipei** is

[<https://www.wikidata.org/wiki/Q1867>](https://www.wikidata.org/wiki/Q1867)

And some example triples (with **IRI prefixes**) look like this:

wd:Q1876 **rdfs:label** "Taipeh"@de

wd:Q1876 **wdt:P17** **wd:Q865**

wd:Q1876 **wdt:P625** "Point(25.03 121.57)"^^**geo:wktLiteral**



Knowledge Graphs 3/8

■ Some widely used knowledge graphs

- **Wikidata:** Started 2013, successor of Freebase (bought by Google in 2010 for 99M \$), crowd-sourced, amazing coverage

18 **B** triples, 52 131 predicates

- **UniProt:** Started 2002, protein sequences, genes, all kinds of metadata, ...

110 **B** triples, 310 predicates

- **PubChem:** Started 2004, chemical compounds, substances, proteins, genes + interrelations

124 **B** triples, 426 predicates





Knowledge Graphs 4/8

■ And more ...

- **OpenStreetMap:** Started 2004, all the geo-data of the world, crowd-sourced, amazing coverage and quality

14 **B** triples, 92 704 predicates

- **DBLP:** Started 1993, publication meta-data for computer science and adjacent fields

0.8 **B** triples, 75 predicates

- Most of the IT-heavy companies maintain their own (very large) knowledge graphs:

Google, Amazon, Microsoft, Meta, Bloomberg, Walmart, Airbnb, ...



■ Historical knowledge graphs

- **Freebase:** Started 2007, acquired by Google in 2010, closed down in 2015

3 B triples, 784 977 predicates



- **YAGO:** Started 2008, derived from Wikipedia info-boxes and WordNet

0.1 B triples, 100 predicates



- **DBpedia:** all kinds of structured data extracted from Wikipedia

0.8 B triples, 54 780 predicates



All three are still mentioned in a lot in papers (because benchmarks were based on them and academia is slow to adapt)



- The standard query language for RDF is
 - SPARQL is a variant of SQL, adapted to the (simple) RDF

```
SELECT ?title ?author ?year WHERE {  
  ?paper dblp:title ?title .  
  ?paper dblp:authoredBy ?author.  
  ?paper dblp:yearOfPublication ?year .  
  FILTER (?year <= 1940)  
}
```

All papers in DBLP
published before
1940

[Query on QLever](#)

- The result of a SPARQL query is always a **table**, in the example:
All assignments to ?title ?author ?year such that the triples exist
in the knowledge graph and the FILTER condition is true
Note: if a paper has **k** authors, there will be **k** rows for it

[all power lines in the EU](#)

■ Interoperability

- One of the great strengths of RDF and SPARQL is the great ease, with which different datasets can be **combined**

For standard databases, this is a nightmare: different and often complex schemas, different identifiers, etc.



In RDF, all you need are additional triples relating the IDs from the two datasets you want to combine



wd:Q183 **wdtn:P402** **osmrel:51477** [in Wikidata]

osmrel:51477 **osm:wikidata** **wd:Q183** [in OpenStreetMap]

wd:Q183

IRI for Germany in Wikidata

osmrel:51477

IRI for Germany in OpenStreetMap

wdtn:P402

predicate for "OpenStreetMap IRI" in Wikidata

osm:wikidata

predicate for "Wikidata IRI" in OpenStreetMap



Knowledge Graphs 8/8

■ A variety of example queries

- Birth places of people with first name X [Wikidata](#)
- Notable events that happened on July 23 [Wikidata](#)
- Proteins with organism and amino sequence [UniProt](#)
- NSAID drugs with small molecular weight [PubChem](#)
- All streets in X [OpenStreetMap](#)
- All countries with official language X [Wikidata+OSM](#)
- Average number of authors per year [DBLP](#)

Side note: These are all running with the same system, with very little configuration per dataset, that is the power of RDF+SPARQL

Finding the right SPARQL query is hard 1/4

■ Example question

- Consider the following simple search request

Oscars of Meryl Streep with corresponding movie

- The result we are looking for is

| | |
|---|-------------------|
| Academy Award for Best Supporting Actress | Kramer vs. Kramer |
| Academy Award for Best Actress | Sophie's Choice |
| Academy Award for Best Actress | The Iron Lady |

- On the next slide, you see the correct SPARQL query on Wikidata



Finding the right SPARQL query is hard 2/4

PREFIX rdfs: <http://www.w3.org/2000/01/rdf-schema#>

PREFIX wdt: <http://www.wikidata.org/prop/direct/>

PREFIX pq: <http://www.wikidata.org/prop/qualifier/>

PREFIX ps: <http://www.wikidata.org/prop/statement/>

PREFIX p: <http://www.wikidata.org/prop/>

PREFIX wd: <http://www.wikidata.org/entity/>

```
SELECT ?movie ?award WHERE {  
  wd:Q873    p:P166    ?statement .  
  ?statement ps:P166   ?award_id .  
  ?statement pq:P1686 ?movie_id .  
  ?award_id  wdt:P31   wd:Q19020 .  
  ?award_id  rdfs:label ?award . FILTER (LANG(?award) = "en")  
  ?movie_id  rdfs:label ?movie . FILTER (LANG(?movie) = "en")  
}
```



Finding the right SPARQL query is hard 3/4

■ What's hard about finding this query?

- Knowing the right prefix definitions

PREFIX rdfs: <<http://www.w3.org/2000/01/rdf-schema#>>

- Knowing the right entity names

wd:Q873 ("Meryl Streep"), wd:Q19020 ("Academy Award")

- Knowing the right predicate names **very hard, even for experts**

p:P166 "won award" from awardee to statement

ps:P166 "won award" from statement to award entity

pq:P1686 "for work" from statement to movie

wdt:P31 "instance of" from "instance" entity to "class" entity

- Knowing the syntax for filtering by language

FILTER (LANG(?award) = "en")



Finding the right SPARQL query is hard 4/4

- A technical solution: query building tools

- **Tool 1:** Wikidata's own query builder

A simple (context-insensitive) autocompletion

Requires **deep expert knowledge** of the prefixes and the right predicate names

- **Tool 2:** QLever's query builder

Context-sensitive suggestions as you type

Requires only relatively **basic knowledge** of RDF/SPARQL, and little or no knowledge about the dataset

Fun fact: the suggestions are themselves computed via SPARQL queries (on the same knowledge graph)

- Ask questions directly in natural language

- Let's try the following question on ChatGPT

Which Oscars did Meryl Streep win and for which movies?

Alas, it works like a charm

So does this mean that we don't need knowledge graphs any longer, but we can just ask our LLM anything?

- Ask questions directly in natural language

- Let's try a slightly more difficult question

First names and how common they are

Older versions of ChatGPT provide wrong answers, newer versions realize that they cannot now

A Large Language Model is no substitute for a database or knowledge graph for at least two reasons:

1. Storing large amounts of information in a LLM is possible, but inefficient (think of the human brain)
2. Some queries require a large "working memory", which LLMs do not have by design

- Use LLMs to help formulate SPARQL queries

- Let's try the following

Give me the SPARQL query on Wikidata for all people with first name X and their birth place with coordinates

ChatGPT gets the query almost right, but the IRI for the first name is wrong

Again, it's hard (and inefficient) for a LLM to store large amounts of relational data, like which entity has which IRI in Wikidata

If we give ChatGPT a hint about the right IRI, it gives us the correct query

■ Definition

- Entity Recognition (ER): identify passages in a text that refer to an entity from a given knowledge graph
- Entity Disambiguation (ED): identify, exactly which entity the passage refers to
- Entity Linking (EL): Entity Recognition + Disambiguation

American athlete Whittington failed to appear in the 2013–14 season due to a torn ACL.

American [[Q30](#)] athlete **Whittington** [[Q21066526](#)] failed to appear in the **2013–14 season** [[Q16192072](#)] due to a **torn ACL** [[Q18912826](#)].

■ Research on Entity Linking

- There are thousands of papers on the problem

[Top venues ER](#)

[Top venues ED](#)

[Top venues EL](#)

- Almost all of these papers claim to improve on previous work and have an evaluation to prove it
- But when you actually try the methods in practice, you frequently make one of the following two experiences:
 1. There is no software or it does not run (anymore)
 2. Lot of hyperparameters and results cannot be replicated
 3. Poor results on datasets not evaluated in the paper

What are the reasons for this? See the following slides

■ Reason 1: Coarse evaluation metrics

- The typical evaluation measures **precision, recall, F1**
- That is fine to get an initial impression of an approach
- But it is very prone to overfitting, especially for benchmarks that have been around a long time
- An absolute must or practically useful entity linking
 1. Look at individual results
 2. Detailed error analysis

■ Reason 2: Benchmarks artifacts and biases

- The following are very frequent in existing benchmarks:

1. Almost exclusive focus on **named entities**

Named entities are almost always capitalized in the English language and hence easy to recognize

2. When going beyond named entities, what is an entity?

No so easy: for example, if everything is an entity that has a Wikipedia article, then almost each piece of text qualifies

3. Many entity mentions are ambiguous

For example, "American" → [Q30](#) or [Q7976](#) or [Q846570](#) ?

4. Many benchmarks have a bias towards types of entites

For example, AIDA-CoNNL: many entities are sports teams

- A general-purpose analysis tool

- **ELEVANT: Entity Linking Evaluation and Analysis Tool**

Let's you examine individual results very nicely (the kind of tool everybody want, but nobody wants to write)

Automatic error analysis of any existing tool

When you develop you own entity linker, it is highly recommended that you use this to inspect your results

- In-depth analysis of existing linkers and benchmarks
 - Compares the best or most well-known linkers in-depth (on existing benchmarks as well as on two new benchmarks)
ReFinED, REL, GENRE, Ambiverse, Neural EL, TagMe, and a simple baseline
 - Points out several problems with several widely used existing benchmarks
AIDA-CoNNL, KORE50, MSNBC, DBpedia Spotlight
 - TODO: what else?



■ Entity Linking and LLMs

- With the right prompt, LLMs can also perform entity linking

Let's try it on ChatGPT (with the GPT-4 model)

- Observations

1. It works surprisingly well right out of the box
2. To obtain great results requires significant fine-tuning
3. In particular, the model often gets the QIDs wrong
4. Entity linking using an LLM is **expensive**

If you want high-quality results with high-performance, tailored solutions are currently still the way to go

References

■ Knowledge Graphs

- <https://query.wikidata.org>
- <https://github.com/ad-freiburg/qllever>
- <https://qllever.cs.uni-freiburg.de>
- <https://aqqv.cs.uni-freiburg.de>

■ Entity Linking

- <https://elevant.cs.uni-freiburg.de>
- <https://arxiv.org/abs/2305.14937>